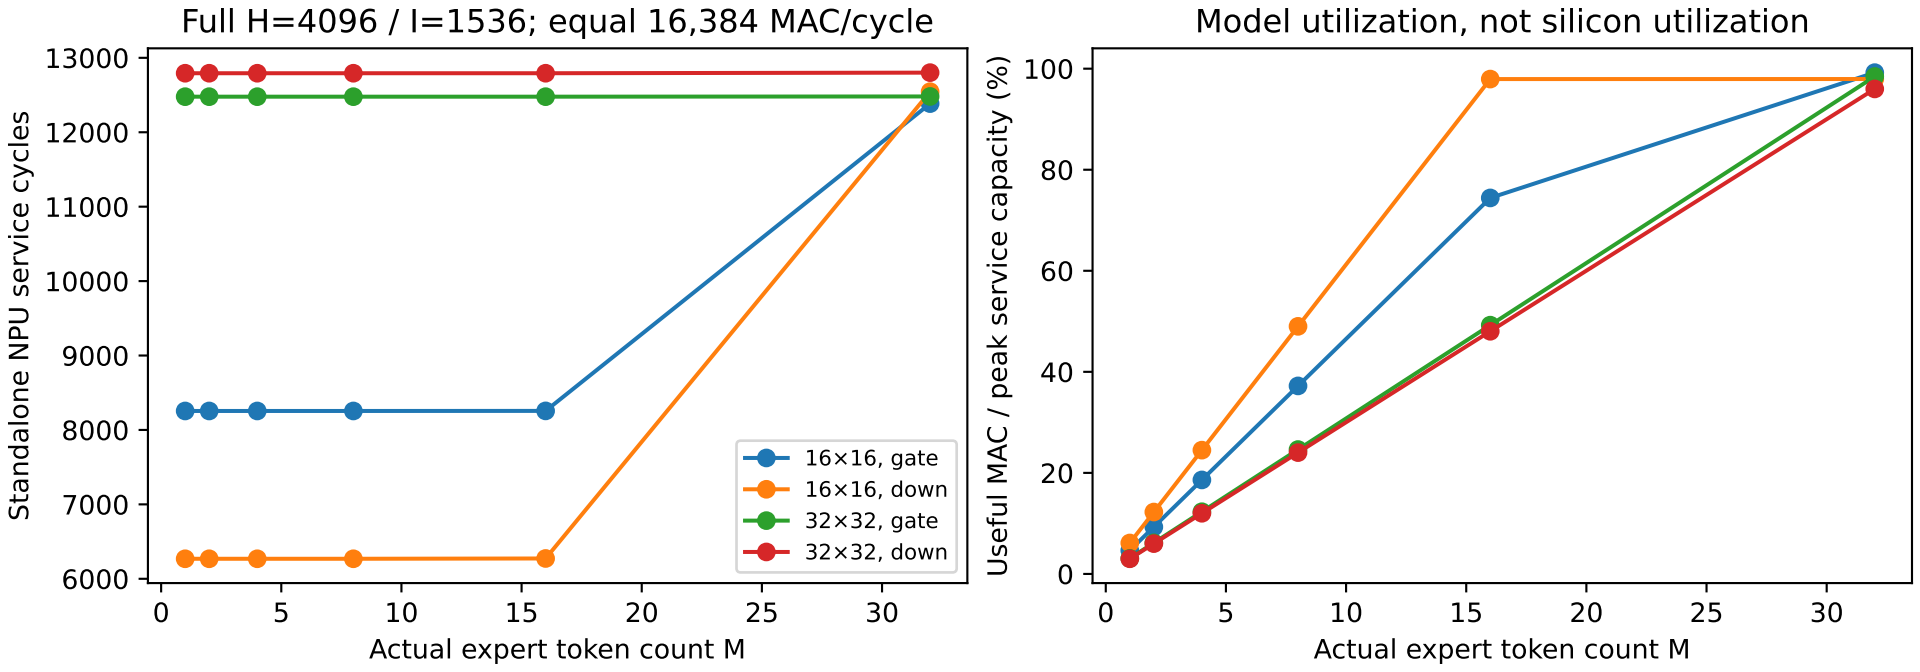
Full H=4096 / I=1536; equal 16,384 MAC/cycle
Standalone NPU service cycles
Actual expert token count M
16×16, gate
16×16, down
32×32, gate
32×32, down
Model utilization, not silicon utilization
Useful MAC / peak service capacity (%)
Actual expert token count M